

1.Package

```
library(e1071)
library(MASS)
library(rpart)
library(tree)
library(randomForest)

## randomForest 4.6-14
## Type rfNews() to see new features/changes/bug fixes.
library(gbm)

## Loaded gbm 2.1.5
library(fastAdaboost)
library(xgboost)
library(ROCR)

## Loading required package: gplots
##
## Attaching package: 'gplots'
## The following object is masked from 'package:stats':
##
##     lowess
library(stringdist)
suppressMessages(library("tidyverse"))
library(tidyverse)
library(caret)

## Loading required package: lattice

## Warning in as.POSIXlt.POSIXct(Sys.time()): unknown timezone 'zone/tz/2019a.
## 1.0/zoneinfo/America/New_York'
##
## Attaching package: 'caret'
## The following object is masked from 'package:purrr':
##
##     lift
```

2. Loading the Dataset

```
location <- read_csv("/Users/effyhou/Desktop/6240mining/hw4/Location.csv")

## Parsed with column specification:
## cols(
##   locationID = col_integer(),
##   regionID = col_integer()
## )
category <- read_csv("/Users/effyhou/Desktop/6240mining/hw4/Category.csv")
```

```
## Parsed with column specification:
## cols(
##   categoryID = col_integer(),
##   parentCategoryID = col_integer()
## )

train <- read_csv("/Users/effyhhou/Desktop/6240mining/hw4/ItemPairs_train.csv")

## Parsed with column specification:
## cols(
##   itemID_1 = col_integer(),
##   itemID_2 = col_integer(),
##   isDuplicate = col_integer(),
##   generationMethod = col_integer()
## )

train_info <- read_csv("/Users/effyhhou/Desktop/6240mining/hw4/ItemInfo_train.csv")

## Parsed with column specification:
## cols(
##   itemID = col_integer(),
##   categoryID = col_integer(),
##   title = col_character(),
##   description = col_character(),
##   images_array = col_character(),
##   attrsJSON = col_character(),
##   price = col_double(),
##   locationID = col_integer(),
##   metroID = col_double(),
##   lat = col_double(),
##   lon = col_double()
## )
```

3. Data pre-processing

```
#First, combine location and regionIDs
train_info <- train_info %>% left_join(location)

## Joining, by = "locationID"

#Second, combine test and train tables with the data in info files

#Some functions to help with the renaming later on
old_cols <- colnames(train)
is_old_column <- function(x){names(x) %in% old_cols}
check_id <- function(x,id="1"){str_sub(names(x),start = -1)==id}
name_adder <- function(x,to_add="1"){paste0(x,to_add)}

#One line dplyr call to combine tables and rename things
train <- train %>%
  left_join(train_info,by=c("itemID_1" = "itemID")) %>%
  rename_if(!is_old_column(.),name_adder,to_add="1") %>%
  left_join(train_info,by=c("itemID_2" = "itemID")) %>%
  rename_if(!is_old_column(.) & !check_id(.,id="1"),name_adder,to_add="2")
```

4. creates features

```
# This function creates features
matchPair <- function(x, y){
  ifelse(is.na(x), ifelse(is.na(y), 3, 2), ifelse(is.na(y), 2, ifelse(x==y, 1, 4)))
}

feature_creator1 <- function(x){
  x %>%
    mutate(#distance
      distance = sqrt((lat1-lat2)^2+(lon1-lon2)^2),
      #same location
      sameLoc=matchPair(locationID1 ,locationID2),
      #same metroID
      samemetro = matchPair(metroID1,metroID2),
      #price
      sameprice=matchPair(price1,price2),
      priceDiff = pmax(price1/price2, price2/price1),
      priceMin = pmin(price1, price2, na.rm=TRUE),
      priceMax = pmax(price1, price2, na.rm=TRUE),
      #title
      titleStringDist = stringdist(title1, title2, method = "jw"),
      titleStringDist2 = (stringdist(title1, title2,
                                     method = "lcs")/pmax(nchar(title1), nchar(title2),
                                                           na.rm=TRUE)),
      titleCharDiff=pmax(nchar(title1)/nchar(title2),
                        nchar(title2)/nchar(title1)),
      titleCharMin = pmin(nchar(title1), nchar(title2), na.rm=TRUE),
      titleCharMax = pmax(nchar(title1), nchar(title2), na.rm=TRUE),
      titleMatch=matchPair(title1,title2),
      descriptionMatch=matchPair(description1,description2),
      descriptionCharDiff = pmax(nchar(description1)/nchar(description2),
                                nchar(description2)/ nchar(description1)),
      descriptionCharMin = pmin( nchar(description1),  nchar(description2), na.rm=TRUE),
      descriptionCharMax = pmax( nchar(description1),  nchar(description2), na.rm=TRUE)
    )
}

feature_creator5 <- function(x){
  x %>%
    mutate(#distance
      distance = sqrt((lat1-lat2)^2+(lon1-lon2)^2),
      #same location
      sameLoc=matchPair(locationID1 ,locationID2),
      #same metroID
      samemetro = matchPair(metroID1,metroID2),
      #price
      sameprice=matchPair(price1,price2),
      priceDiff = pmax(price1/price2, price2/price1),
      priceMin = pmin(price1, price2, na.rm=TRUE),
      priceMax = pmax(price1, price2, na.rm=TRUE),
```

```

#title
titleStringDist = stringdist(title1, title2, method = "jw"),
titleStringDist2 = (stringdist(title1, title2,
                              method = "lcs")/pmax(nchar(title1), nchar(title2),
                                                    na.rm=TRUE)),

titleCharDiff=pmax(nchar(title1)/nchar(title2),
                  nchar(title2)/nchar(title1)),
titleCharMin = pmin(nchar(title1), nchar(title2), na.rm=TRUE),
titleCharMax = pmax(nchar(title1), nchar(title2), na.rm=TRUE),
titleMatch=matchPair(title1,title2),
descriptionMatch=matchPair(description1,description2),
descriptionCharDiff = pmax(nchar(description1)/nchar(description2),
                          nchar(description2)/ nchar(description1)),
descriptionCharMin = pmin( nchar(description1),  nchar(description2), na.rm=TRUE),
descriptionCharMax = pmax( nchar(description1),  nchar(description2), na.rm=TRUE),

# title-discription distance
title_discription_Dist_jw_1_1 = stringdist(title1, description1, method = "jw"),
title_discription_Dist_jw_2_2 = stringdist(title2, description2, method = "jw"),
title_discription_Dist_ja_1_1 = stringdist(title1, description1, method = "jaccard"),
title_discription_Dist_ja_2_2 = stringdist(title2, description2, method = "jaccard"),
title_discription_Dist_co_1_1 = stringdist(title1, description1, method = "cosine"),
title_discription_Dist_co_2_2 = stringdist(title2, description2, method = "cosine")

)
}

train1<- train%>% feature_creator1
train5<- train%>% feature_creator5

train1[is.na(train1)] <- -9999
train1[train1==Inf] <- -9999
train5[is.na(train5)] <- -9999
train5[train5==Inf] <- -9999

```

5. Randomly subsample and split

sample train with feature1

```

set.seed(123)
subtrain1 <- sample_frac(train1,0.03)

subtrain1 <- subtrain1 %>% mutate(isDuplicate=factor(isDuplicate))
subtrain1 <- subtrain1 %>% select(isDuplicate,distance:descriptionCharMax)

#Split the data into train, test and validation
spec1<- c(sample_train1 = 1/3, sample_test1 = 1/3, sample_valid1 = 1/3)

split1 <- sample(cut(
  seq(nrow(subtrain1)),
  nrow(subtrain1)*cumsum(c(0,spec1)),
  labels = names(spec1)

```

```

))

res1 <- split(subtrain1, split1)

sample_train1 <- res1$sample_train1
sample_test1 <- res1$sample_test1
sample_valid1 <- res1$sample_valid1

```

sample train with feature5

```

set.seed(123)
subtrain5 <- sample_frac(train5,0.03)

subtrain5 <- subtrain5 %>% mutate(isDuplicate=factor(isDuplicate))
subtrain5 <- subtrain5 %>% select(isDuplicate,distance:descriptionCharMax)

#Split the data into train, test and validation
spec5<- c(sample_train5 = 1/3, sample_test5 = 1/3, sample_valid5= 1/3)

split5 <- sample(cut(
  seq(nrow(subtrain5)),
  nrow(subtrain5)*cumsum(c(0,spec5)),
  labels = names(spec5)
))

res5 <- split(subtrain5, split5)

sample_train5 <- res5$sample_train5
sample_test5 <- res5$sample_test5
sample_valid5 <- res5$sample_valid5

```

6. Fit 10 different models on the training data

Based on HW4 I choose top 5 models : XGboost, RandomForest,logistic, LDA, gbm

6.1 h2o randomForest

```

#h2o.shutdown()
library(h2o)

##
## -----
##
## Your next step is to start H2O:
##   > h2o.init()
##
## For H2O package documentation, ask for help:
##   > ??h2o
##
## After starting H2O, you can use the Web UI at http://localhost:54321

```



```

|
|=====| 100%
features1<-colnames(sample_train1)[!(colnames(sample_train1) %in% c("isDuplicate"))]
validationHex1<-as.h2o(sample_valid1)

```

```

##
|
|                                     | 0%
|
|=====| 100%

```

```
testHex1<-as.h2o(sample_test1)
```

```

##
|
|                                     | 0%
|
|=====| 100%

```

```

rf1 <- h2o.randomForest(x=features1,
                        y="isDuplicate",
                        training_frame = sample_trainHex1,
                        validation_frame = validationHex1,
                        ntree=500,
                        seed = 123)

```

```

##
|
|                                     | 0%
|
|                                     | 1%
|=
|                                     | 2%
|=
|                                     | 2%
|==
|                                     | 5%
|===
|                                     | 8%
|====
|                                     | 11%
|=====
|                                     | 14%
|=====
|                                     | 17%
|=====
|                                     | 19%
|=====
|                                     | 22%
|=====
|                                     | 25%
|=====
|                                     | 28%
|=====
|                                     | 30%

```

```

|
|=====| 34%
|=====| 37%
|=====| 41%
|=====| 43%
|=====| 44%
|=====| 47%
|=====| 51%
|=====| 54%
|=====| 58%
|=====| 61%
|=====| 64%
|=====| 68%
|=====| 71%
|=====| 75%
|=====| 78%
|=====| 82%
|=====| 85%
|=====| 88%
|=====| 90%
|=====| 93%
|=====| 96%
|=====| 100%

```

```

#predict validation
rf1_pred_valid<-predict(rf1 ,validationHex1, probability=TRUE)[3]

```

```

##
|
| 0%
|=====| 100%

```



```

rf1_pred_valid<-as.vector(rf1_pred_valid)
#predict test
rf1_pred_test<-predict(rf1 ,testHex1, probability=TRUE)[3]

##
|
|
|
|=====| 100%

rf1_pred_test<-as.vector(rf1_pred_test)

sample_trainHex5<-as.h2o(sample_train5)

##
|
|
|
|=====| 100%

features5<-colnames(sample_train5)[!(colnames(sample_train5) %in% c("isDuplicate"))]
validationHex5<-as.h2o(sample_valid5)

##
|
|
|
|=====| 100%

testHex5<-as.h2o(sample_test5)

##
|
|
|
|=====| 100%

rf5<- h2o.randomForest(x=features5,
                       y="isDuplicate",
                       training_frame = sample_trainHex5,
                       validation_frame = validationHex5,
                       ntree=500,
                       seed = 123)

##
|
|
|
|=
|
|=
|
|==
|
|===
|
|=====| 6%

```

=====	9%
=====	12%
=====	15%
=====	19%
=====	22%
=====	25%
=====	29%
=====	32%
=====	34%
=====	37%
=====	38%
=====	41%
=====	42%
=====	45%
=====	48%
=====	51%
=====	55%
=====	58%
=====	61%
=====	65%
=====	68%
=====	71%
=====	75%
=====	78%
=====	82%
=====	85%
=====	88%

```

|=====| 92%
|
|=====| 95%
|
|=====| 98%
|
|=====| 100%

#predict validation
rf5_pred_valid<-predict(rf5 ,validationHex5, probability=TRUE)[3]

##
|
|
|=====| 100%

rf5_pred_valid<-as.vector(rf5_pred_valid)
#predict test
rf5_pred_test<-predict( rf5 ,testHex5, probability=TRUE)[3]

##
|
|
|=====| 100%

rf5_pred_test<-as.vector(rf5_pred_test)

```

6.2 xgboost

```

maxTrees <- 200
shrinkage <- 0.10
gamma <- 1
depth <- 10
minChildWeight <- 40
colSample <- 0.85
subSample <- 0.85
earlyStopRound <- 4

xg1_features<-colnames(sample_train1)[!(colnames(sample_train1) %in% c("isDuplicate"))]
d_train1 <- xgb.DMatrix(as.matrix(sample_train1[, xg1_features]), label=as.numeric(sample_train1$isDupl
d_validation1 <- sample_valid1%>%
  select(-isDuplicate) %>%
  as.matrix %>%
  xgb.DMatrix(label=as.numeric(sample_valid1$isDuplicate)-1)

test_p1<-sample_test1[, -1]
d_test1 <- xgb.DMatrix(as.matrix(test_p1))

xgb1 <- xgboost(params=list(max_depth=depth,
                             eta=shrinkage,
                             gamma=gamma,
                             colsample_bytree=colSample,

```

```
min_child_weight=minChildWeight),  
data=d_train1,  
nrounds=90,  
objective="binary:logistic",  
eval_metric="auc")    #0.855005
```

```
## [1] train-auc:0.803136  
## [2] train-auc:0.806879  
## [3] train-auc:0.809378  
## [4] train-auc:0.811136  
## [5] train-auc:0.813328  
## [6] train-auc:0.815260  
## [7] train-auc:0.816865  
## [8] train-auc:0.818603  
## [9] train-auc:0.819839  
## [10] train-auc:0.820743  
## [11] train-auc:0.821799  
## [12] train-auc:0.822579  
## [13] train-auc:0.823784  
## [14] train-auc:0.824586  
## [15] train-auc:0.825443  
## [16] train-auc:0.826205  
## [17] train-auc:0.827532  
## [18] train-auc:0.828335  
## [19] train-auc:0.828973  
## [20] train-auc:0.829890  
## [21] train-auc:0.830706  
## [22] train-auc:0.831196  
## [23] train-auc:0.832043  
## [24] train-auc:0.832697  
## [25] train-auc:0.833395  
## [26] train-auc:0.833942  
## [27] train-auc:0.834635  
## [28] train-auc:0.835169  
## [29] train-auc:0.835903  
## [30] train-auc:0.836743  
## [31] train-auc:0.837548  
## [32] train-auc:0.837838  
## [33] train-auc:0.838305  
## [34] train-auc:0.838727  
## [35] train-auc:0.839320  
## [36] train-auc:0.839842  
## [37] train-auc:0.840496  
## [38] train-auc:0.841006  
## [39] train-auc:0.841181  
## [40] train-auc:0.841626  
## [41] train-auc:0.841827  
## [42] train-auc:0.842337  
## [43] train-auc:0.842954  
## [44] train-auc:0.843337  
## [45] train-auc:0.843778  
## [46] train-auc:0.844016  
## [47] train-auc:0.844326  
## [48] train-auc:0.844626
```

```
## [49] train-auc:0.844923
## [50] train-auc:0.845365
## [51] train-auc:0.845763
## [52] train-auc:0.846170
## [53] train-auc:0.846396
## [54] train-auc:0.846778
## [55] train-auc:0.847135
## [56] train-auc:0.847619
## [57] train-auc:0.847864
## [58] train-auc:0.848027
## [59] train-auc:0.848134
## [60] train-auc:0.848624
## [61] train-auc:0.848689
## [62] train-auc:0.848759
## [63] train-auc:0.849083
## [64] train-auc:0.849215
## [65] train-auc:0.849543
## [66] train-auc:0.849721
## [67] train-auc:0.849901
## [68] train-auc:0.850234
## [69] train-auc:0.850339
## [70] train-auc:0.850559
## [71] train-auc:0.850680
## [72] train-auc:0.851008
## [73] train-auc:0.851209
## [74] train-auc:0.851365
## [75] train-auc:0.851530
## [76] train-auc:0.851910
## [77] train-auc:0.852200
## [78] train-auc:0.852293
## [79] train-auc:0.852704
## [80] train-auc:0.853051
## [81] train-auc:0.853203
## [82] train-auc:0.853479
## [83] train-auc:0.853646
## [84] train-auc:0.853780
## [85] train-auc:0.853833
## [86] train-auc:0.853877
## [87] train-auc:0.853976
## [88] train-auc:0.854362
## [89] train-auc:0.854731
## [90] train-auc:0.855005
```

```
#predict validation:
```

```
xgb1_pred_valid <- predict( xgb1, d_validation1)
```

```
#predict test:
```

```
xgb1_pred_test <- predict( xgb1, d_test1)
```

```
xg5_features<-colnames(sample_train5)[!(colnames(sample_train5) %in% c("isDuplicate"))]
```

```
d_train5 <- xgb.DMatrix(as.matrix(sample_train5[, xg5_features]), label=as.numeric(sample_train5$isDupl.
```

```
d_validation5 <- sample_valid5%>%
```

```
  select(-isDuplicate) %>%
```

```

as.matrix %>%
xgb.DMatrix(label=as.numeric(sample_valid5$isDuplicate)-1)

test_p5<-sample_test5[,-1]
d_test5 <- xgb.DMatrix(as.matrix(test_p5))

xgb5 <- xgboost(params=list(max_depth=depth,
                             eta=shrinkage,
                             gamma=gamma,
                             colsample_bytree=colSample,
                             min_child_weight=minChildWeight),
               data=d_train5,
               nrounds=100,
               objective="binary:logistic",
               eval_metric="auc") #0.857057

## [1] train-auc:0.803136
## [2] train-auc:0.806879
## [3] train-auc:0.809378
## [4] train-auc:0.811136
## [5] train-auc:0.813328
## [6] train-auc:0.815260
## [7] train-auc:0.816865
## [8] train-auc:0.818603
## [9] train-auc:0.819839
## [10] train-auc:0.820743
## [11] train-auc:0.821799
## [12] train-auc:0.822579
## [13] train-auc:0.823784
## [14] train-auc:0.824586
## [15] train-auc:0.825443
## [16] train-auc:0.826205
## [17] train-auc:0.827532
## [18] train-auc:0.828335
## [19] train-auc:0.828973
## [20] train-auc:0.829890
## [21] train-auc:0.830706
## [22] train-auc:0.831196
## [23] train-auc:0.832043
## [24] train-auc:0.832697
## [25] train-auc:0.833395
## [26] train-auc:0.833942
## [27] train-auc:0.834635
## [28] train-auc:0.835169
## [29] train-auc:0.835903
## [30] train-auc:0.836743
## [31] train-auc:0.837548
## [32] train-auc:0.837838
## [33] train-auc:0.838305
## [34] train-auc:0.838727
## [35] train-auc:0.839320
## [36] train-auc:0.839842
## [37] train-auc:0.840496

```

```
## [38] train-auc:0.841006
## [39] train-auc:0.841181
## [40] train-auc:0.841626
## [41] train-auc:0.841827
## [42] train-auc:0.842337
## [43] train-auc:0.842954
## [44] train-auc:0.843337
## [45] train-auc:0.843778
## [46] train-auc:0.844016
## [47] train-auc:0.844326
## [48] train-auc:0.844626
## [49] train-auc:0.844923
## [50] train-auc:0.845365
## [51] train-auc:0.845763
## [52] train-auc:0.846170
## [53] train-auc:0.846396
## [54] train-auc:0.846778
## [55] train-auc:0.847135
## [56] train-auc:0.847619
## [57] train-auc:0.847864
## [58] train-auc:0.848027
## [59] train-auc:0.848134
## [60] train-auc:0.848624
## [61] train-auc:0.848689
## [62] train-auc:0.848759
## [63] train-auc:0.849083
## [64] train-auc:0.849215
## [65] train-auc:0.849543
## [66] train-auc:0.849721
## [67] train-auc:0.849901
## [68] train-auc:0.850234
## [69] train-auc:0.850339
## [70] train-auc:0.850559
## [71] train-auc:0.850680
## [72] train-auc:0.851008
## [73] train-auc:0.851209
## [74] train-auc:0.851365
## [75] train-auc:0.851530
## [76] train-auc:0.851910
## [77] train-auc:0.852200
## [78] train-auc:0.852293
## [79] train-auc:0.852704
## [80] train-auc:0.853051
## [81] train-auc:0.853203
## [82] train-auc:0.853479
## [83] train-auc:0.853646
## [84] train-auc:0.853780
## [85] train-auc:0.853833
## [86] train-auc:0.853877
## [87] train-auc:0.853976
## [88] train-auc:0.854362
## [89] train-auc:0.854731
## [90] train-auc:0.855005
## [91] train-auc:0.855339
```

```
## [92] train-auc:0.855377
## [93] train-auc:0.855482
## [94] train-auc:0.855837
## [95] train-auc:0.856207
## [96] train-auc:0.856639
## [97] train-auc:0.856724
## [98] train-auc:0.856912
## [99] train-auc:0.856941
## [100] train-auc:0.857057

#predict validation:
xgb5_pred_valid <- predict( xgb5, d_validation5)

#predict test:
xgb5_pred_test <- predict( xgb5, d_test5)

xgb5_prediction<- prediction(xgb5_pred_test,labels=sample_test1$isDuplicate)

performance(xgb5_prediction,"auc")@y.values[[1]] # 0.8178197

## [1] 0.8178197
```

6.3 LDA

```
#LDA:

lda1<- lda(isDuplicate~.,data = sample_train1)
lda1

## Call:
## lda(isDuplicate ~ ., data = sample_train1)
##
## Prior probabilities of groups:
##      0      1
## 0.5770877 0.4229123
##
## Group means:
##      distance sameLoc samemetro sameprice priceDiff priceMin priceMax
## 0 0.9250093 1.188380 2.293170 3.280890 376367.8016 127854.5 20985361
## 1 2.0356004 1.248044 2.553316 2.422259 -977.3178 7301817.6 14364313
## titleStringDist titleStringDist2 titleCharDiff titleCharMin titleCharMax
## 0 0.3212639 0.7500445 1.44943 23.34200 30.02908
## 1 0.1849947 0.4133287 1.36139 21.29207 25.96435
## titleMatch descriptionMatch descriptionCharDiff descriptionCharMin
## 0 3.610207 3.862596 1.812630 286.8211
## 1 2.896372 3.744447 1.199938 214.5414
## descriptionCharMax
## 0 420.1962
## 1 284.4412
##
## Coefficients of linear discriminants:
## LD1
## distance 1.299021e-02
```



```
## sameLoc          8.816746e-02
## samemetro        2.214140e-01
## sameprice        -4.277468e-01
## priceDiff        -2.078212e-10
## priceMin         3.609131e-11
## priceMax         -7.811239e-12
## titleStringDist   4.990981e-01
## titleStringDist2 -1.357229e+00
## titleCharDiff     6.566496e-02
## titleCharMin     -5.814435e-03
## titleCharMax     -3.338133e-03
## titleMatch        -1.373988e-01
## descriptionMatch  -1.059220e-01
## descriptionCharDiff -6.615096e-06
## descriptionCharMin -7.360458e-05
## descriptionCharMax -3.820759e-04
```

```
#predict validation
lda1_pred_valid<- lda1 %>%
  predict(sample_valid1) %>%
  (function(x) x$posterior[,2])
lda1_pred_valid <-as.vector(lda1_pred_valid)
```

```
#predict test
lda1_pred_test <- lda1 %>%
  predict(sample_test1) %>%
  (function(x) x$posterior[,2])
lda1_pred_test<-as.vector(lda1_pred_test)
```

```
####Feature 5
lda5<- lda(isDuplicate~.,data = sample_train5)
lda5
```

```
## Call:
## lda(isDuplicate ~ ., data = sample_train5)
##
## Prior probabilities of groups:
##      0      1
## 0.5770877 0.4229123
##
## Group means:
##      distance sameLoc samemetro sameprice priceDiff priceMin priceMax
## 0 0.9250093 1.188380 2.293170 3.280890 376367.8016 127854.5 20985361
## 1 2.0356004 1.248044 2.553316 2.422259 -977.3178 7301817.6 14364313
## titleStringDist titleStringDist2 titleCharDiff titleCharMin titleCharMax
## 0 0.3212639 0.7500445 1.44943 23.34200 30.02908
## 1 0.1849947 0.4133287 1.36139 21.29207 25.96435
## titleMatch descriptionMatch descriptionCharDiff descriptionCharMin
## 0 3.610207 3.862596 1.812630 286.8211
## 1 2.896372 3.744447 1.199938 214.5414
## descriptionCharMax
## 0 420.1962
## 1 284.4412
##
## Coefficients of linear discriminants:
```

```

##                                LD1
## distance                      1.299021e-02
## sameLoc                       8.816746e-02
## samemetro                     2.214140e-01
## sameprice                     -4.277468e-01
## priceDiff                     -2.078212e-10
## priceMin                      3.609131e-11
## priceMax                      -7.811239e-12
## titleStringDist               4.990981e-01
## titleStringDist2              -1.357229e+00
## titleCharDiff                 6.566496e-02
## titleCharMin                  -5.814435e-03
## titleCharMax                  -3.338133e-03
## titleMatch                    -1.373988e-01
## descriptionMatch              -1.059220e-01
## descriptionCharDiff           -6.615096e-06
## descriptionCharMin            -7.360458e-05
## descriptionCharMax            -3.820759e-04

#predict validation
lda5_pred_valid<- lda5 %>%
  predict(sample_valid5) %>%
  (function(x) x$posterior[,2])
lda5_pred_valid <-as.vector(lda5_pred_valid)

#predict test
lda5_pred_test <- lda5 %>%
  predict(sample_test5) %>%
  (function(x) x$posterior[,2])
lda5_pred_test<-as.vector(lda5_pred_test)

```

6.4 GBM

```
sample_trainHex1<-as.h2o(sample_train1)
```

```

##
|
|
|
|=====| 100%

```

```
validationHex1<-as.h2o(sample_valid1)
```

```

##
|
|
|
|=====| 100%

```

```
testHex1<-as.h2o(sample_test1)
```

```

##
|
|
|
|=====| 100%

```



```
##
|
|
| 0%
|=====| 100%
```

```
gbm_pred_test1<-as.vector(gbm_pred_test1)
```

```
sample_trainHex5<-as.h2o(sample_train5)
```

```
##
|
|
| 0%
|=====| 100%
```

```
validationHex5<-as.h2o(sample_valid5)
```

```
##
|
|
| 0%
|=====| 100%
```

```
testHex5<-as.h2o(sample_test5)
```

```
##
|
|
| 0%
|=====| 100%
```

```
gbm5 <- h2o.gbm(
  ## standard model parameters
  x = features5,
  y="isDuplicate",
  training_frame = sample_trainHex5,
  validation_frame = validationHex5,
  ntrees = 500,
  learn_rate=0.07,
  sample_rate = 0.8,
  col_sample_rate = 0.6,
  seed = 1234,
  max_depth=7
)
```

```
##
|
|
| 0%
|==| 3%
|====| 6%
|=====| 10%
```

```

|
|=====| 13%
|=====| 23%
|=====| 34%
|=====| 45%
|=====| 55%
|=====| 88%
|=====| 100%

```

```
#predict validation
```

```
gbm_pred_valid5<-predict(gbm5 ,validationHex5, probability=TRUE)[3]
```

```
##
```

```

|
|
|
|=====| 100%

```

```
gbm_pred_valid5<-as.vector(gbm_pred_valid5)
```

```
#predict test
```

```
gbm_pred_test5<-predict(gbm5,testHex5, probability=TRUE)[3]
```

```
##
```

```

|
|
|
|=====| 100%

```

```
gbm_pred_test5<-as.vector(gbm_pred_test5)
```

6.5 Logistic regression

```
lg1 <- glm(isDuplicate ~ .,data=sample_train1,family="binomial")
```

```
## Warning: glm.fit: algorithm did not converge
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```

lg1_pred_valid <- lg1 %>%
  predict(sample_valid1,type="response")
lg1_pred_valid<-as.vector(lg1_pred_valid)

```

```

lg1_pred_test <- lg1 %>%
  predict(sample_test1,type="response")
lg1_pred_test<-as.vector(lg1_pred_test)

```

```
lg5 <- glm(isDuplicate ~ .,data=sample_train5,family="binomial")
```

```
## Warning: glm.fit: algorithm did not converge
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
lg5_pred_valid <- lg5 %>%  
  predict(sample_valid5,type="response")  
lg5_pred_valid<-as.vector(lg5_pred_valid)
```

```
lg5_pred_test <- lg5 %>%  
  predict(sample_test5,type="response")  
lg5_pred_test<-as.vector(lg5_pred_test)
```

7. Stacking

7.1 stacking model (stacking 10 different models)

Use 10 Models:2 logistic regression, 2 xgboost, 2 gbm, 2 LDA, 2 RandomForest;

Use feature 1 and feature 5

```
stack_v<-cbind(lda1_pred_valid,lda5_pred_valid,  
              gbm_pred_valid1,gbm_pred_valid5,  
              lg1_pred_valid,lg5_pred_valid,  
              xgb1_pred_valid,xgb5_pred_valid,  
              rf1_pred_valid,rf5_pred_valid)  
  
stack_t<-cbind(lda1_pred_test,lda5_pred_test,  
              gbm_pred_test1,gbm_pred_test5,  
              lg1_pred_test,lg5_pred_test,  
              xgb1_pred_test,xgb5_pred_test,  
              rf1_pred_test,rf5_pred_test)  
  
stack_v_xg<- xgb.DMatrix(as.matrix(stack_v), label=as.numeric(sample_valid1$isDuplicate)-1)  
  
modelStack<- xgboost(params=list(max_depth=depth,  
                                eta=shrinkage,  
                                gamma=gamma,  
                                colsample_bytree=colSample,  
                                min_child_weight=minChildWeight),  
                    data=stack_v_xg,  
                    nrounds=100,  
                    objective="binary:logistic",  
                    eval_metric="auc") #0.843325
```

```
## [1] train-auc:0.824238  
## [2] train-auc:0.826193  
## [3] train-auc:0.826974  
## [4] train-auc:0.827889  
## [5] train-auc:0.828571  
## [6] train-auc:0.829523  
## [7] train-auc:0.830061  
## [8] train-auc:0.830501  
## [9] train-auc:0.830977
```

```
## [10] train-auc:0.831344
## [11] train-auc:0.831721
## [12] train-auc:0.832048
## [13] train-auc:0.832310
## [14] train-auc:0.832662
## [15] train-auc:0.832851
## [16] train-auc:0.833139
## [17] train-auc:0.833596
## [18] train-auc:0.833935
## [19] train-auc:0.834204
## [20] train-auc:0.834476
## [21] train-auc:0.834782
## [22] train-auc:0.835042
## [23] train-auc:0.835327
## [24] train-auc:0.835538
## [25] train-auc:0.835863
## [26] train-auc:0.836041
## [27] train-auc:0.836282
## [28] train-auc:0.836601
## [29] train-auc:0.836863
## [30] train-auc:0.837138
## [31] train-auc:0.837264
## [32] train-auc:0.837545
## [33] train-auc:0.837644
## [34] train-auc:0.837820
## [35] train-auc:0.837970
## [36] train-auc:0.838293
## [37] train-auc:0.838405
## [38] train-auc:0.838578
## [39] train-auc:0.838649
## [40] train-auc:0.838882
## [41] train-auc:0.839007
## [42] train-auc:0.839161
## [43] train-auc:0.839283
## [44] train-auc:0.839496
## [45] train-auc:0.839654
## [46] train-auc:0.839755
## [47] train-auc:0.839872
## [48] train-auc:0.840046
## [49] train-auc:0.840142
## [50] train-auc:0.840184
## [51] train-auc:0.840257
## [52] train-auc:0.840441
## [53] train-auc:0.840505
## [54] train-auc:0.840634
## [55] train-auc:0.840698
## [56] train-auc:0.840748
## [57] train-auc:0.840861
## [58] train-auc:0.840905
## [59] train-auc:0.840951
## [60] train-auc:0.841007
## [61] train-auc:0.841047
## [62] train-auc:0.841103
## [63] train-auc:0.841256
```

```
## [64] train-auc:0.841336
## [65] train-auc:0.841361
## [66] train-auc:0.841387
## [67] train-auc:0.841439
## [68] train-auc:0.841690
## [69] train-auc:0.841719
## [70] train-auc:0.841825
## [71] train-auc:0.841847
## [72] train-auc:0.841875
## [73] train-auc:0.841998
## [74] train-auc:0.842034
## [75] train-auc:0.842057
## [76] train-auc:0.842085
## [77] train-auc:0.842123
## [78] train-auc:0.842148
## [79] train-auc:0.842184
## [80] train-auc:0.842212
## [81] train-auc:0.842262
## [82] train-auc:0.842294
## [83] train-auc:0.842313
## [84] train-auc:0.842379
## [85] train-auc:0.842676
## [86] train-auc:0.842692
## [87] train-auc:0.842714
## [88] train-auc:0.842786
## [89] train-auc:0.842856
## [90] train-auc:0.842908
## [91] train-auc:0.842930
## [92] train-auc:0.842969
## [93] train-auc:0.843163
## [94] train-auc:0.843181
## [95] train-auc:0.843212
## [96] train-auc:0.843240
## [97] train-auc:0.843260
## [98] train-auc:0.843275
## [99] train-auc:0.843300
## [100] train-auc:0.843325
```

```
modelStack_predict<- predict(modelStack,stack_t)
```

```
#AUC
```

```
stack_prediction<- prediction(modelStack_predict,labels=sample_test1$isDuplicate)
```

```
performance(stack_prediction,"auc")@y.values[[1]] # 0.819282
```

```
## [1] 0.819282
```

Stacking model has higher AUC (0.819282) than the single XGboost model (which has the highest AUC score (0.8178197) among 10 different models).

7.2 Use Stacking model to obtain classifications

```
class_stack<-ifelse (modelStack_predict > 0.5,1,0)
```